

Data@Anz- Data-Analysis

August 18, 2021

0.0.1 1.0 Import Libraries

```
[12]: #remove warnings
import warnings
warnings.filterwarnings('ignore')

# Data analysis and wrangling
import pandas as pd
import numpy as np
import statistics
# for data visualization
import seaborn as sns
%matplotlib inline
from matplotlib import pyplot as plt
from matplotlib import style
get_ipython().run_line_magic('matplotlib', 'inline')
# For Dates Conversion
import datetime
# for selection of Algorithms
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn import linear_model
from sklearn.linear_model import Perceptron
from sklearn.linear_model import SGDClassifier
from sklearn.model_selection import train_test_split

[13]: from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import cross_val_score
# Import more libraries
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.metrics import accuracy_score, roc_auc_score, confusion_matrix, _
    classification_report
```

```

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LinearRegression
from statsmodels.tools.eval_measures import rmse
from sklearn import linear_model
from sklearn import metrics
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error as MSE
pd.set_option('max_rows', None)

```

0.0.2 2.0 Getting and Loading Dataset

```

[16]: #Getting the dataset
import os

# Path of the file to read
Anz_dataset = pd.read_csv(r"C:\Users\User\Downloads\DATA SCIENCE\anz.csv")

```

0.0.3 3.0 Exploratory Analysis

```

[17]: #checking columns
Anz_dataset.columns

```

```

[17]: Index(['status', 'card_present_flag', 'bpay_biller_code', 'account',
            'currency', 'long_lat', 'txn_description', 'merchant_id',
            'merchant_code', 'first_name', 'balance', 'date', 'gender', 'age',
            'merchant_suburb', 'merchant_state', 'extraction', 'amount',
            'transaction_id', 'country', 'customer_id', 'merchant_long_lat',
            'movement'],
            dtype='object')

```

```

[18]: #checking info about the Anz_dataset
Anz_dataset.head()

print(Anz_dataset.info())

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12043 entries, 0 to 12042
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status                12043 non-null object
1   card_present_flag     7717 non-null  float64
2   bpay_biller_code      885 non-null   object
3   account               12043 non-null object
4   currency              12043 non-null object
5   long_lat              12043 non-null object

```

6	txn_description	12043 non-null	object
7	merchant_id	7717 non-null	object
8	merchant_code	883 non-null	float64
9	first_name	12043 non-null	object
10	balance	12043 non-null	float64
11	date	12043 non-null	object
12	gender	12043 non-null	object
13	age	12043 non-null	int64
14	merchant_suburb	7717 non-null	object
15	merchant_state	7717 non-null	object
16	extraction	12043 non-null	object
17	amount	12043 non-null	float64
18	transaction_id	12043 non-null	object
19	country	12043 non-null	object
20	customer_id	12043 non-null	object
21	merchant_long_lat	7717 non-null	object
22	movement	12043 non-null	object

dtypes: float64(4), int64(1), object(18)

memory usage: 2.1+ MB

None

There missing values in some columns

```
[19]: #printing data shape
print('Anz_dataset shape: ',Anz_dataset.shape)
#printing the number of rows
print('Number of rows: ',len(Anz_dataset))
```

Anz_dataset shape: (12043, 23)

Number of rows: 12043

```
[20]: # checking statistics summary of Anz_dataset
Anz_dataset.describe()
```

```
[20]:
```

	card_present_flag	merchant_code	balance	age \
count	7717.000000	883.0	12043.000000	12043.000000
mean	0.802644	0.0	14704.195553	30.582330
std	0.398029	0.0	31503.722652	10.046343
min	0.000000	0.0	0.240000	18.000000
25%	1.000000	0.0	3158.585000	22.000000
50%	1.000000	0.0	6432.010000	28.000000
75%	1.000000	0.0	12465.945000	38.000000
max	1.000000	0.0	267128.520000	78.000000

	amount
count	12043.000000
mean	187.933588
std	592.599934

```

min      0.100000
25%     16.000000
50%     29.000000
75%     53.655000
max     8835.980000

```

Observation:

There is inconsistency and lots of zero values in merchant code column. It's probably a categorical column.

4.1 Checking for unique customers

```

[21]: # Checking for the 100 unique customers
print("Number of unique customer ID's: ", Anz_dataset.customer_id.
      ↪unique())
print("Number of unique transaction ID's: ", Anz_dataset.transaction_id.
      ↪unique())
print("Number of unique accounts: ", Anz_dataset.account.nunique())

```

```

Number of unique customer ID's: 100
Number of unique transaction ID's: 12043
Number of unique accounts: 100

```

4.2 Checking the format of Date

```

[22]: Anz_dataset.date.describe()
      #Anz_dataset.date.count()

```

```

[22]: count      12043
      unique        91
      top      9/28/2018
      freq        174
      Name: date, dtype: object

```

Observation:

One day is missing. Date format is consistent

0.0.4 5.0 Missing Values

```

[23]: #Checking for missing values


def missing_data(Anz_dataset, title):
    fig, ax = plt.subplots(figsize=(5,5))
    plt.title(title)
    sns.heatmap(Anz_dataset.isnull(), cbar=False)

```

```

[24]: missing_data(Anz_dataset, "Anz_dataset")

```



Observation:

The missing columns are evident from the graph

0.0.5 5.1 Checking Percentage of missing values

```
[25]: # checking percentage of missing values
missing_value = Anz_dataset.isnull().sum()
missing_value = missing_value[missing_value > 0]
percentage_missing_value = round(missing_value / len(Anz_dataset), 3) * 100

pd.DataFrame({"Number of missing_value": missing_value, "Percentage": _
↪percentage_missing_value}).sort_values(by = "Percentage", ascending = False)
```

[25]:		Number of missing_value	Percentage
	bpay_biller_code	11158	92.7
	merchant_code	11160	92.7
	card_present_flag	4326	35.9
	merchant_id	4326	35.9
	merchant_suburb	4326	35.9
	merchant_state	4326	35.9
	merchant_long_lat	4326	35.9

Observation:

merchant_code and bpay_biller_code have a high percentage of missing values. We are dropping the two columns due to high percentage of missing values.

0.0.6 5.2 Treating Null values

```
[26]: Anz_dataset.drop(columns=['merchant_code', 'bpay_biller_code'], inplace=True)
```

```
[27]: #dropping null values in merchant
df = pd.DataFrame(Anz_dataset.merchant_state)
new_df = df.dropna()
new_df.head()
```

```
[27]: merchant_state
0      QLD
1      NSW
2      NSW
3      QLD
4      QLD
```

Confirming the drop

```
[28]: missing_value = new_df.isnull().sum()
print(missing_value)
#Anz_dataset.head(2)
```

```
merchant_state    0
dtype: int64
```

```
[29]: Anz_dataset.columns
```

```
[29]: Index(['status', 'card_present_flag', 'account', 'currency', 'long_lat',
        'txn_description', 'merchant_id', 'first_name', 'balance', 'date',
        'gender', 'age', 'merchant_suburb', 'merchant_state', 'extraction',
        'amount', 'transaction_id', 'country', 'customer_id',
        'merchant_long_lat', 'movement'],
        dtype='object')
```

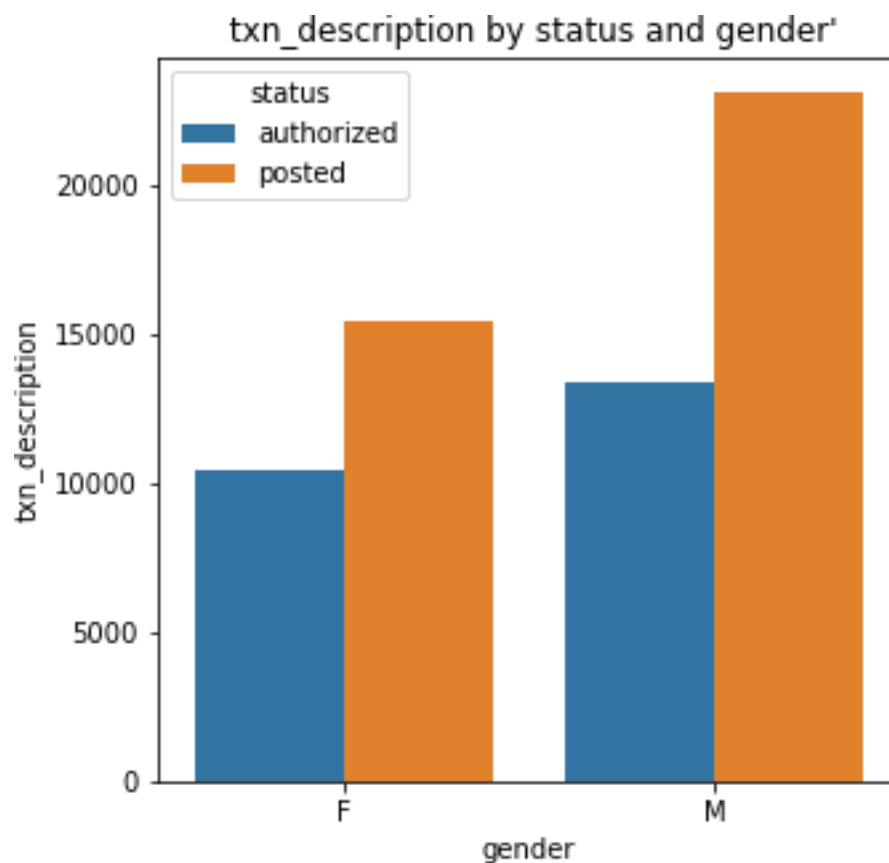
0.1 6.0 Analysis of Features

Distribution plot function

0.1.1 Analysis of gender and status

```
[30]: def bar_chart_compare(dataset, feature1, feature2=None, title = "↵  
↵txn_description by status and gender"):↵  
    plt.figure(figsize = [5,5])↵  
    plt.title(title)↵  
    g = sns.barplot(x=feature1, y="balance", hue=feature2, ci=None,↵  
    ↵data=dataset).set_ylabel("txn_description")
```

```
[31]: bar_chart_compare(Anz_dataset, "gender", "status")
```




There were more males authorized and posted than females

0.1.2 Gender and Amount

```
[32]: plt.figure(figsize = (10,5))  
sns.distplot(Anz_dataset.amount)  
plt.title("amount")
```

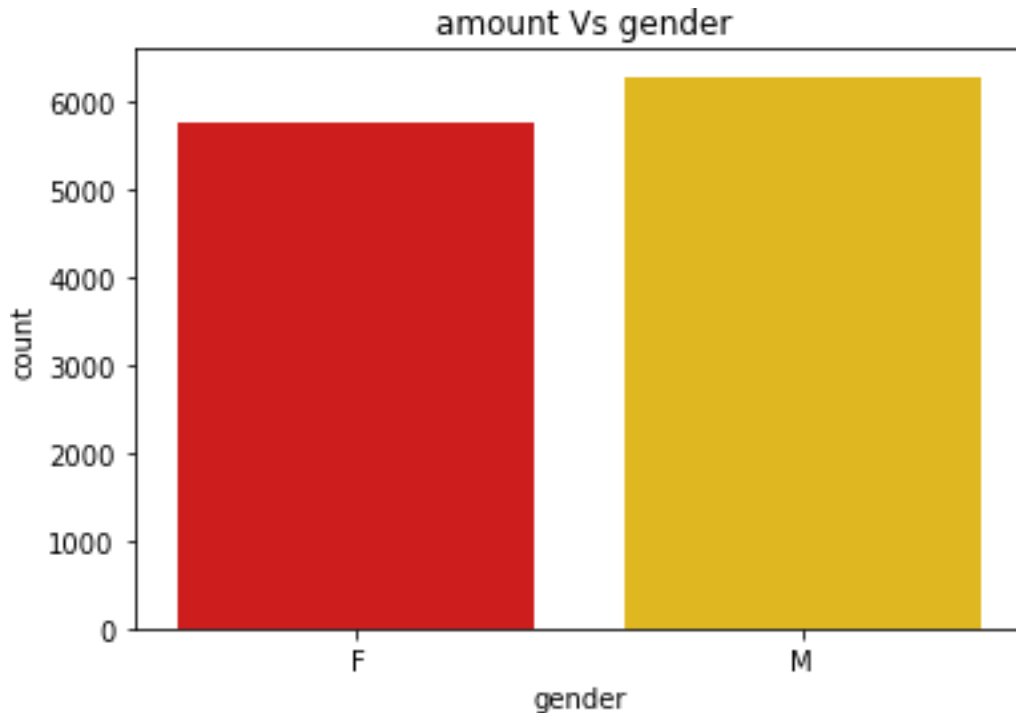
```
[32]: Text(0.5, 1.0, 'amount')
```



amount distribution is right or positively skewed

```
[33]: sns.countplot(x= 'gender',data = Anz_dataset, palette ="hot")  
plt.title(' amount Vs gender')
```

```
[33]: Text(0.5, 1.0, ' amount Vs gender')
```

Observation:

More males engaged in transaction than females. the distribution

Transaction

```
[34]: # checking Average
tran_average = statistics.mean(Anz_dataset.amount)
print('average transactional amount is:', round(tran_average, 2))
```

average transactional amount is: 187.93

```
[35]: Anz_dataset.groupby('txn_description').amount.mean()
```

```
[35]: txn_description
INTER BANK      86.699461
PAY/SALARY     1898.728029
PAYMENT         77.613077
PHONE BANK     106.099010
POS             40.407412
SALES-POS       39.909789
Name: amount, dtype: float64
```

```
[36]: # Statistical summary
Anz_dataset.groupby('txn_description').amount.describe()
```

```
[36]:
```

txn_description	count	mean	std	min	25%	50% \
INTER BANK	742.0	86.699461	198.706044	16.0	26.000	39.000
PAY/SALARY	883.0	1898.728029	1150.364621	576.0	1013.670	1626.480
PAYMENT	2600.0	77.613077	152.310315	15.0	32.000	42.500
PHONE BANK	101.0	106.099010	245.999695	21.0	36.000	43.000
POS	3783.0	40.407412	165.771678	0.1	12.035	19.430
SALES-POS	3934.0	39.909789	132.734185	0.1	12.160	20.035

txn_description	75%	max
INTER BANK	83.000	1956.00
PAY/SALARY	2538.680	8835.98
PAYMENT	70.000	1981.00
PHONE BANK	67.000	1916.00
POS	33.155	7081.09
SALES-POS	34.575	4233.00


Observation:

There is a huge difference between the minimum and the maximum numbers. There is high confidence interval, it means the sample mean is not reliable.

0.1.3 Analysizing Features gender and Number of transaction

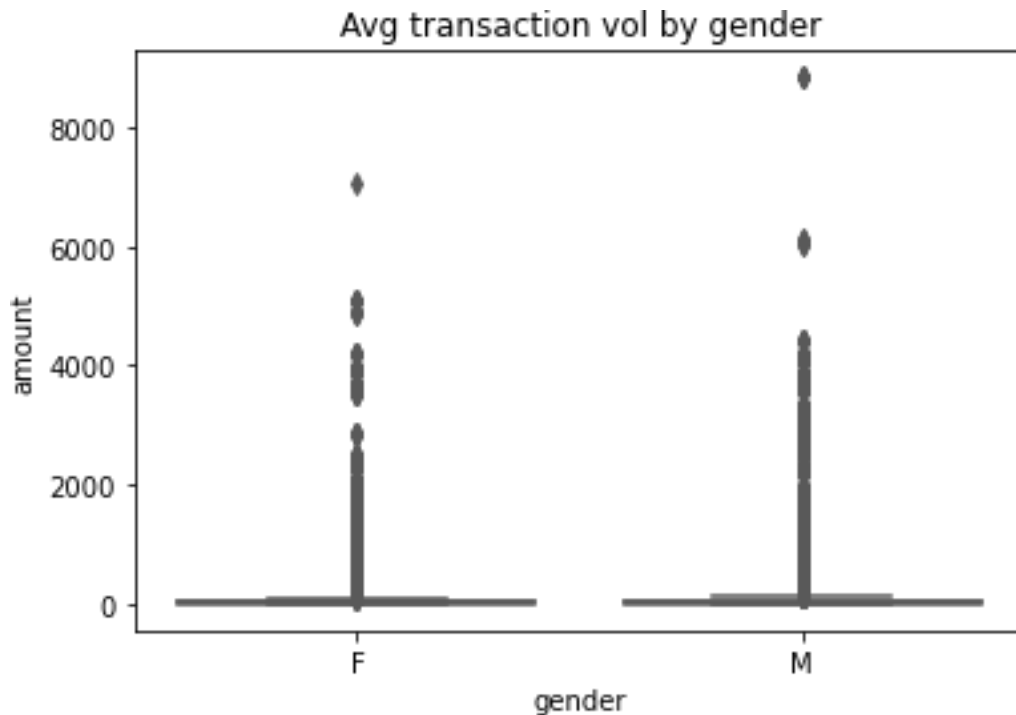
```
[37]: sns.countplot(x= "gender",data = Anz_dataset, palette ="hls")
plt.title(" Number of transaction Vs gender")
```

```
[37]: Text(0.5, 1.0, ' Number of transaction Vs gender')
```



```
[38]: sns.boxplot(x= "gender", y= "amount", data = Anz_dataset, palette = "Set2")  
plt.title("Avg transaction vol by gender")
```

```
[38]: Text(0.5, 1.0, 'Avg transaction vol by gender')
```



There are more male customers with transactions than females This is a confirmed case.

[39]: #Average transaction volume by state and movement

```
Anz_dataset.merchant_suburb.dropna().head()
```

```
[39]: 0      Ashmore
      1      Sydney
      2      Sydney
      3      Buderim
      4  Mermaid Beach
      Name: merchant_suburb, dtype: object
```

[40]: Anz_dataset.card_present_flag.dropna().isnull().sum()


[40]: 0

Observation:

PaySalary is the highest transaction, the margin is wide.

```
[41]: plt.figure(figsize = (10,5))
      sns.barplot(x="txn_description",y= "amount",data = Anz_dataset,
      ↪ hue="movement",palette = "tab10")
      plt.title("Avg transaction volume by category and movement")
```

[41]: Text(0.5, 1.0, 'Avg transaction volume by category and movement')




Observations:

salaries have the highest transaction and were paid with credit. others used debit and no credit. There is high confidence interval, implying the sample mean was not reliable as an estimate of the true amount of the salary, Interbank and Phone bank. This means the average portrayed is false.


```
[42]: def plot_swarm_movement(dataset, feature1, feature2, title, size = (15,5)):
    fig, ax = plt.subplots(figsize=(18,5))
    # Turns off grid on the left Axis.
    ax.grid(True)
    plt.xticks(list(range(0,100,2)))
    sns.swarmplot(y=feature1, x=feature2, hue='movement', data=Anz_dataset).
    ↪ set_title(title)
```

```
[43]: plot_swarm_movement(Anz_dataset, 'balance', 'movement', 'movement Swarmplot for_
    ↪ movement and balance')
```



From the swarm plot the debit is performing well than the credit.


```
[44]: plot_swarm_movement(Anz_dataset, 'amount', 'merchant_state', 'movement Swarmplot_↵for amount and merchant_state')
```



ACT's average transaction volume is the highest but it is the state with the lowest number of transaction. This means that, the company needs to focus on ACT since its average transaction volume is high. While the NSW and VIC have a high number of transactions, their average transaction volume is relatively low. Hence, little effort should be put there.

```
[45]: sns.countplot(x='merchant_state',data = Anz_dataset, palette ="rocket")  
plt.title('Number transaction volume by state')
```


```
[45]: Text(0.5, 1.0, 'Number transaction volume by state')
```



Observation: There are more male customers with transactions than females ACT's average transaction volume is the highest but it is the state with the lowest number of transaction

```
[46]: sns.countplot(x= "movement",data = Anz_dataset,palette = "mako")  
      plt.title("transaction Vs movement")
```

```
[46]: Text(0.5, 1.0, 'transaction Vs movement')
```



There are more debit transactions than credit


```
[ ]: plt.figure(figsize = (10,5))
sns.barplot(x="merchant_state", y="amount", data = Anz_dataset, 
            hue="movement", palette = "coolwarm")
plt.title("Average transaction volume by state and movement")
```

Observation: ACT has the highest average transaction volume but the variance is quite large.

0.1.4 Age distribution

```
[48]: plt.figure(figsize = (10,5))
sns.distplot(Anz_dataset.age)
plt.title("Age_Distribution")
```

```
[48]: Text(0.5, 1.0, 'Age_Distribution')
```

lowest transactions came from people after 50 year while most of transactions came from people in 20's.


Transaction Volume

```
[49]: daily = pd.DataFrame(Anz_dataset.groupby("date").amount.sum())
      daily.head()
```

```
[49]:      amount
      date
10/1/2018  40823.03
10/10/2018 29399.50
10/11/2018 20851.67
10/12/2018 40658.20
10/13/2018 10140.81
```

```
[50]: fig, ax = plt.subplots(figsize = (12, 5))
      ax.plot(daily.index, daily.amount)
      plt.title("Transaction volume from 1/8/2018 to 31/10/2018")
      plt.xlabel("Date")
      plt.ylabel("Daily Transaction volume")
```

```
[50]: Text(0, 0.5, 'Daily Transaction volume')
```




There is a similar pattern of rising and dropping transaction volume

0.1.5 Balance Distribution


```
[51]: plt.figure(figsize = (12, 10))
      sns.boxplot(Anz_dataset.balance)
```

```
[51]: <AxesSubplot:xlabel='balance'>
```



```
[52]: plt.figure(figsize = (12, 5))  
sns.distplot(Anz_dataset.balance)  
plt.title("Balance distribution")
```

```
[52]: Text(0.5, 1.0, 'Balance distribution')
```



The balance session is rightly skewed, hence, needs to be corrected before the model.

```
[53]: customer_monthly_volume = pd.DataFrame(Anz_dataset.groupby('customer_id').
      ↪amount.
      sum()/3)
      customer_monthly_volume.head()
```

```
[53]:          amount
customer_id
CUS-1005756958  5422.990000
CUS-1117979751  11328.123333
CUS-1140341822   5670.200000
CUS-1147642491   9660.273333
CUS-1196156254  12016.906667
```

```
[54]: pd.DataFrame(Anz_dataset.groupby('customer_id').amount.sum())
```


```
[54]:          amount
customer_id
CUS-1005756958  16268.97
CUS-1117979751  33984.37
CUS-1140341822  17010.60
CUS-1147642491  28980.82
CUS-1196156254  36050.72
CUS-1220154422  20596.11
CUS-1233833708  10385.54
CUS-1271030853  29079.78
CUS-127297539   21856.81
CUS-134193016   17230.81
CUS-134833760   32243.17
```

CUS-1388323263	18680.22
CUS-1433879684	15563.36
CUS-1462656821	34511.20
CUS-1478398256	27698.18
CUS-1499065773	19609.82
CUS-1604596597	19221.00
CUS-1609060617	23378.97
CUS-1614226872	19286.44
CUS-1617121891	33085.75
CUS-164374203	27722.98
CUS-1646183815	10845.25
CUS-1646621553	22141.51
CUS-1654129794	10587.42
CUS-1669695324	23070.56
CUS-1739931018	10652.72
CUS-1790886359	28489.54
CUS-1816693151	40215.54
CUS-1842679196	12438.05
CUS-1892177589	23222.91
CUS-1896554896	20818.80
CUS-1928710999	22729.12
CUS-2031327464	35832.97
CUS-2059096722	26234.46
CUS-2083971310	21230.42
CUS-2110742437	13645.24
CUS-2142601169	23696.45
CUS-2155701614	37943.79
CUS-2178051368	16033.50
CUS-2206365095	15062.14
CUS-2283904812	16942.84
CUS-2317998716	21422.69
CUS-2348881191	18754.06
CUS-2370108457	20006.47
CUS-2376382098	32198.92
CUS-2484453271	12630.60
CUS-2487424745	26211.59
CUS-2500783281	26408.22
CUS-2505971401	29711.92
CUS-2599279756	14763.14
CUS-261674136	36786.13
CUS-2630892467	13614.83
CUS-2650223890	17635.02
CUS-2663907001	33459.56
CUS-2688605418	20550.24
CUS-2695611575	21479.97
CUS-2738291516	45409.16
CUS-2819545904	28265.48

CUS-2977593493	13522.71
CUS-3026014945	29074.34
CUS-3117610635	22820.90
CUS-3129499595	21006.53
CUS-3142625864	42688.30
CUS-3151318058	16424.00
CUS-3174332735	24950.27
CUS-3180318393	22792.27
CUS-3201519139	14535.11
CUS-3249305314	32401.50
CUS-325142416	27483.01
CUS-3255104878	13259.72
CUS-326006476	15687.97
CUS-331942311	14922.51
CUS-3325710106	20515.56
CUS-3336454548	34020.50
CUS-3378712515	22761.96
CUS-3395687666	14216.01
CUS-3431016847	15324.86
CUS-3462882033	22801.59
CUS-3702001629	15463.59
CUS-3716701010	23233.50
CUS-3904958894	14692.24
CUS-3989008654	19160.97
CUS-4023861240	26521.55
CUS-4123612273	22312.12
CUS-4142663097	36588.25
CUS-423725039	13864.12
CUS-443776336	16932.26
CUS-495599312	21502.62
CUS-497688347	15400.63
CUS-511326734	16615.46
CUS-51506836	24100.75
CUS-527400765	36543.61
CUS-537508723	23516.20
CUS-55310383	28367.16
CUS-586638664	17222.50
CUS-72755508	11438.37
CUS-809013380	18810.09
CUS-860700529	18099.88
CUS-880898248	11462.45
CUS-883482547	36639.41


```
[55]: plt.figure(figsize = (10, 5))  
sns.distplot(customer_monthly_volume.amount)  
plt.title("Customers transaction volume monthly")
```

[55]: Text(0.5, 1.0, 'Customers transaction volume monthly')



```
[56]: plt.figure(figsize = (12, 10))  
sns.boxplot(Anz_dataset.amount)
```

[56]: <AxesSubplot:xlabel='amount'>



The amount session is also rightly skewed and hence has to be transformed before the model.

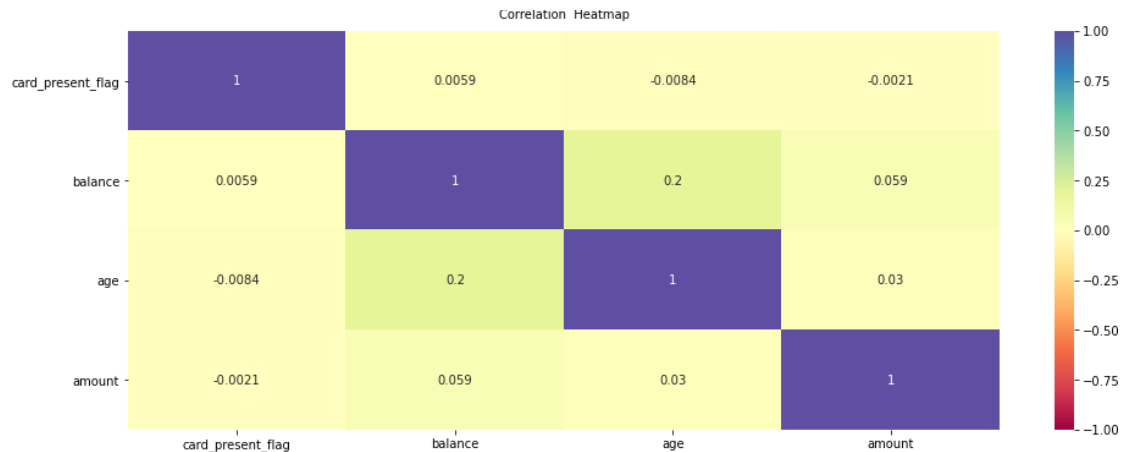
0.1.6 Checking Correlation

```
[57]: Anz_dataset.corr()
```

```
[57]:
```

	card_present_flag	balance	age	amount
card_present_flag	1.000000	0.005925	-0.008405	-0.002074
balance	0.005925	1.000000	0.199329	0.059178
age	-0.008405	0.199329	1.000000	0.029980
amount	-0.002074	0.059178	0.029980	1.000000

```
[58]: plt.figure(figsize=(16, 6))
heatmap = sns.heatmap(Anz_dataset.corr(), vmin=-1, vmax=1, annot=True,
cmap="Spectral")
heatmap.set_title('Correlation Heatmap',fontdict={'fontsize':10}, pad=10);
```

There are correlation in the dataset, We do not need to drop any column.

```
[59]: Anz_dataset.head(1)
      #type(Anz_dataset)
```

```
[59]:      status  card_present_flag      account currency      long_lat \
0  authorized                1.0  ACC-1598451071      AUD  153.41 -27.95

      txn_description      merchant_id first_name  balance \
0      POS 81c48296-73be-44a7-befa-d053f48ce7cd      Diana  35.39

      date  ... age  merchant_suburb  merchant_state \
0  8/1/2018  ...  26      Ashmore      QLD

      extraction amount      transaction_id \
0  2018-08-01T01:01:15.000+0000  16.25  a623070bfead4541a6b0fff8a09e706c

      country  customer_id  merchant_long_lat  movement
0  Australia  CUS-2487424745  153.38 -27.99  debit

[1 rows x 21 columns]
```

```
[ ]:
```

1 Data@ANZ Programme Task2

1.1 Predictive Analytics Task

1.1.1 Reindexing

```
[60]: #print(Anz_dataset['status'])
```

```
[61]: #Anz_data=Anz_dataset.set_index('status')
```

1.1.2 Creating more features

1.1.3 Feature Annual Salary

```
[62]: #setting the background for seaborn
sns.set_style('ticks')
#exclude all except pay/salary in txn_description
Salary = Anz_dataset[(Anz_dataset['txn_description'] == 'PAY/SALARY')]

#summing amount by customer id to sum up all the salary during the 3_
↳months Salary=
Salary.pivot_table(index='customer_id',values='amount',aggfunc=np.sum)
Salary.reset_index(inplace=True)
#creating annual salary by multiplying the sum of amount by 4(the
↳data is 3 months data and so a year which is 12 months need to be_
↳multiplied by 4)
Salary['Annual_Salary'] = (round(Salary['amount']*4,2))
Salary.head()
```

```
[62]:
```

	index	status	card_present_flag	account	currency	long_lat	\
0	50	posted	NaN	ACC-588564840	AUD	151.27 -33.76	
1	61	posted	NaN	ACC-1650504218	AUD	145.01 -37.93	
2	64	posted	NaN	ACC-3326339947	AUD	151.18 -33.80	
3	68	posted	NaN	ACC-3541460373	AUD	145.00 -37.83	
4	70	posted	NaN	ACC-2776252858	AUD	144.95 -37.76	

	txn_description	merchant_id	first_name	balance	...	merchant_suburb	\
0	PAY/SALARY	NaN	Isaiah	8342.11	...	NaN	
1	PAY/SALARY	NaN	Marissa	2040.58	...	NaN	
2	PAY/SALARY	NaN	Eric	3158.51	...	NaN	
3	PAY/SALARY	NaN	Jeffrey	2517.66	...	NaN	
4	PAY/SALARY	NaN	Kristin	2271.79	...	NaN	

	merchant_state	extraction	amount	\
0	NaN	2018-08-01T11:00:00.000+0000	3903.95	
1	NaN	2018-08-01T12:00:00.000+0000	1626.48	
2	NaN	2018-08-01T12:00:00.000+0000	983.36	
3	NaN	2018-08-01T13:00:00.000+0000	1408.08	
4	NaN	2018-08-01T13:00:00.000+0000	1068.04	

	transaction_id	country	customer_id	\
0	9ca281650e5d482d9e53f85e959baa66	Australia	CUS-1462656821	

```

1 1822eb0e1bbe4c9e95ebbb0fa2cc4323 Australia CUS-2500783281
2 bd62b1799a454cedbbb56364f7c40cbf Australia CUS-326006476
3 0d95c7c932bb48e5b44c2637bdd3efe9 Australia CUS-1433879684
4 f50ccf1195214d14a0acbfcb5a265193 Australia CUS-4123612273

```

```

merchant_long_lat movement Annual_Salary
0 NaN credit 15615.80
1 NaN credit 6505.92
2 NaN credit 3933.44
3 NaN credit 5632.32
4 NaN credit 4272.16

```

[5 rows x 23 columns]

1.1.4 Feature Annual Balance

```

[63]: #exclude all except pay/salary in txn_description
Balance = Anz_dataset[(Anz_dataset["txn_description"] == "PAY/SALARY")]
#summing amount by customer id to sum up all the balance during the
↳ 3 months
Balance= Balance.pivot_table(index="customer_id",values="balance",aggfunc=np.
↳ sum)
Balance.reset_index(inplace=True)

```

```

[64]: #creating annual balance by multiplying the sum of amount by 4(the
↳ data is 3_months data and so a year which is 12 months need to be
↳ multiplied by 4)
Balance["Annual_balance"] = (round(Balance["balance"]*4,2))
Balance.head()

```

```

[64]: customer_id balance Annual_balance
0 CUS-1005756958 61342.65 245370.60
1 CUS-1117979751 83700.42 334801.68
2 CUS-1140341822 35050.32 140201.28
3 CUS-1147642491 114575.08 458300.32
4 CUS-1196156254 166920.02 667680.08

```

1.1.5 Feature Expenses

```

[65]: #exclude credit in movement.
Expense = Anz_dataset[Anz_dataset["movement"] == "debit"]
#summing amount by customer id to sum up the expenses during the 3 months
Expense = Expense.pivot_table(index="customer_id",values="amount",aggfunc=np.
↳ sum)
#creating annual expenses by multiplying the sum of amount by 4(the
↳ data is 3_months data and so a year which is 12 months need to be
↳ multiplied by 4)

```

```
Expense["Annual_Expense"] = round(Expense["amount"]*4)
Expense.reset_index(inplace=True)
Expense.head()
```


```
[65]:
```

	customer_id	amount	Annual_Expense
0	CUS-1005756958	3652.86	14611.0
1	CUS-1117979751	8933.82	35735.0
2	CUS-1140341822	5511.54	22046.0
3	CUS-1147642491	6732.75	26931.0
4	CUS-1196156254	8724.61	34898.0

1.2 Distribution Plots

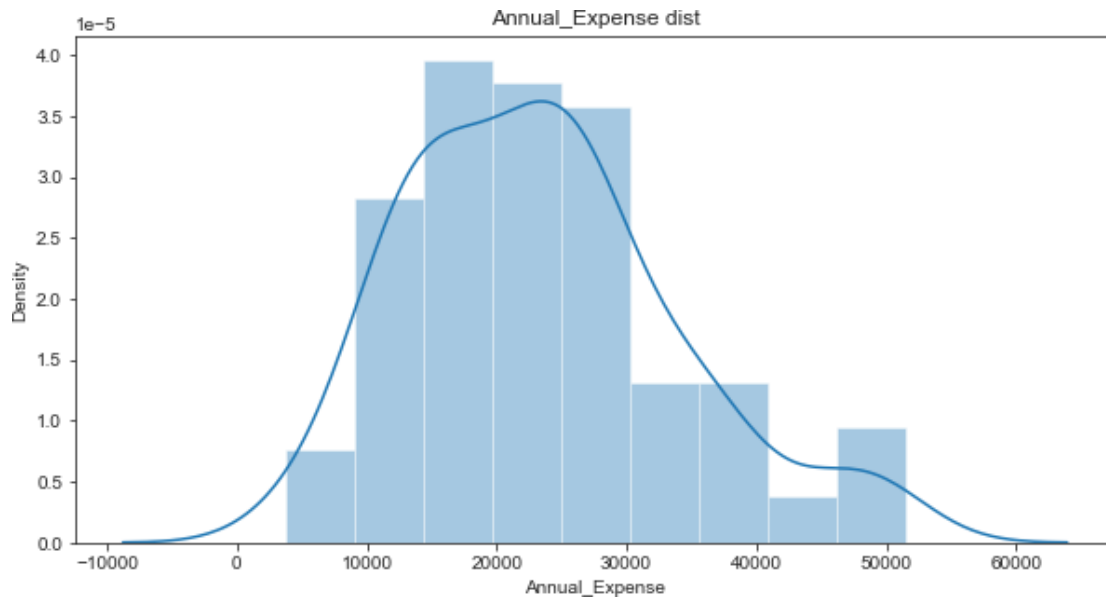
```
[66]: plt.figure(figsize = (12,5))
sns.distplot(Balance.Annual_balance)
plt.title("Annual_balance dist")
```

```
[66]: Text(0.5, 1.0, 'Annual_balance dist')
```



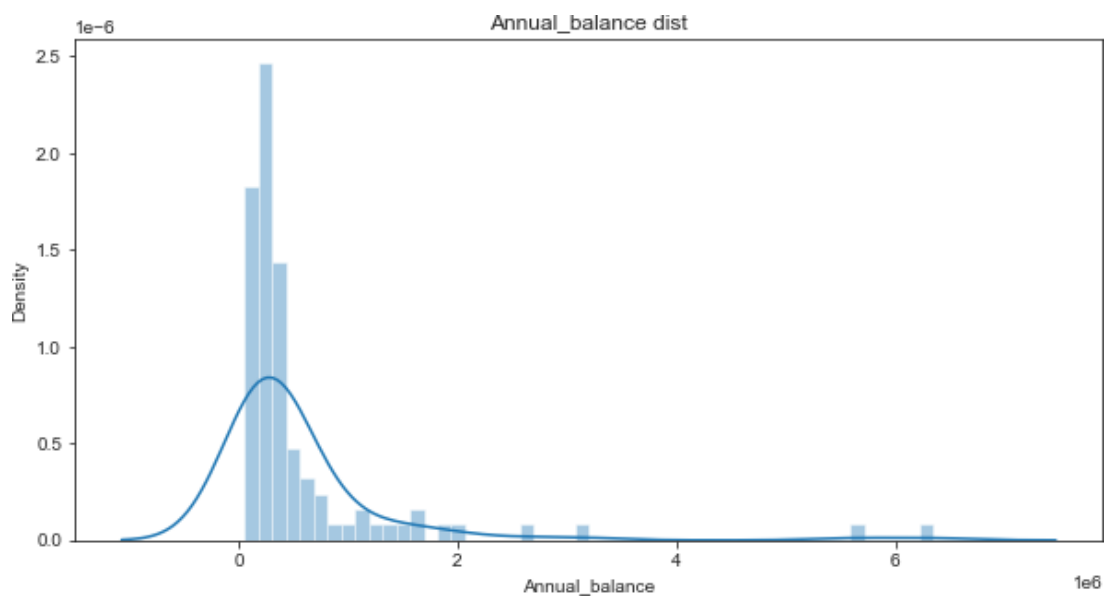
```
[67]: plt.figure(figsize = (10,5))
sns.distplot(Expense.Annual_Expense)
plt.title("Annual_Expense dist")
```

```
[67]: Text(0.5, 1.0, 'Annual_Expense dist')
```



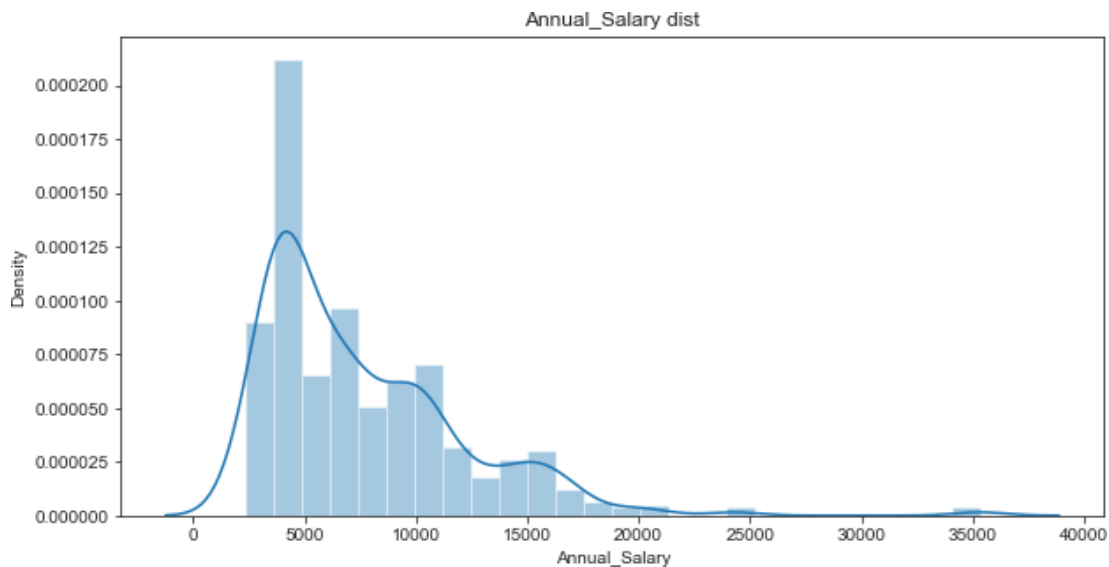
```
[68]: plt.figure(figsize = (10,5))
sns.distplot(Balance.Annual_balance)
plt.title("Annual_balance dist")
```

```
[68]: Text(0.5, 1.0, 'Annual_balance dist')
```



```
[69]: plt.figure(figsize = (10,5))
sns.distplot(Salary.Annual_Salary)
plt.title("Annual_Salary dist")
```

```
[69]: Text(0.5, 1.0, 'Annual_Salary dist')
```



Merging age and gender with new feaures created

```
[70]: Salary=pd.
↳merge(Salary,Anz_dataset[["customer_id","age","gender"]],on="customer_id",how="left")
Salary.drop_duplicates(inplace=True)
Salary.head()
```

```
[70]:
```

	index	status	card_present_flag	account	currency	long_lat
0	50	posted	NaN	ACC-588564840	AUD	151.27 -33.76
116	61	posted	NaN	ACC-1650504218	AUD	145.01 -37.93
200	64	posted	NaN	ACC-3326339947	AUD	151.18 -33.80
247	68	posted	NaN	ACC-3541460373	AUD	145.00 -37.83
402	70	posted	NaN	ACC-2776252858	AUD	144.95 -37.76

	txn_description	merchant_id	first_name	balance
0	PAY/SALARY	NaN	Isaiah	8342.11
116	PAY/SALARY	NaN	Marissa	2040.58
200	PAY/SALARY	NaN	Eric	3158.51
247	PAY/SALARY	NaN	Jeffrey	2517.66
402	PAY/SALARY	NaN	Kristin	2271.79

extraction	amount	transaction_id
------------	--------	----------------

```

0    2018-08-01T11:00:00.000+0000    3903.95    9ca281650e5d482d9e53f85e959baa66
116  2018-08-01T12:00:00.000+0000    1626.48    1822eb0e1bbe4c9e95ebbb0fa2cc4323
200  2018-08-01T12:00:00.000+0000    983.36    bd62b1799a454cedbbb56364f7c40cbf
247  2018-08-01T13:00:00.000+0000    1408.08    0d95c7c932bb48e5b44c2637bdd3efe9
402  2018-08-01T13:00:00.000+0000    1068.04    f50ccf1195214d14a0acbfc5a265193

```

```

      country    customer_id merchant_long_lat    movement    Annual_Salary \
0    Australia    CUS-1462656821          NaN    credit      15615.80
116  Australia    CUS-2500783281          NaN    credit      6505.92
200  Australia    CUS-326006476          NaN    credit      3933.44
247  Australia    CUS-1433879684          NaN    credit      5632.32
402  Australia    CUS-4123612273          NaN    credit      4272.16

```

```

      age_y    gender_y
0         23         M
116        23         F
200        22         M
247        24         M
402        43         F

```

[5 rows x 25 columns]

1.2.1 Merging balance data with salary data

```

[71]: Salary = pd.merge(Salary,Expense,on='customer_id',how='left')
Salary.drop_duplicates(inplace=True)
Salary.head()

```

```

[71]:   index    status    card_present_flag    account    currency    long_lat \
0      50    posted          NaN    ACC-588564840      AUD    151.27 -33.76
1      61    posted          NaN    ACC-1650504218      AUD    145.01 -37.93
2      64    posted          NaN    ACC-3326339947      AUD    151.18 -33.80
3      68    posted          NaN    ACC-3541460373      AUD    145.00 -37.83
4      70    posted          NaN    ACC-2776252858      AUD    144.95 -37.76

```

```

      txn_description    merchant_id    first_name    balance    ... \
0      PAY/SALARY          NaN    Isaiah    8342.11    ...
1      PAY/SALARY          NaN    Marissa    2040.58    ...
2      PAY/SALARY          NaN    Eric    3158.51    ...
3      PAY/SALARY          NaN    Jeffrey    2517.66    ...
4      PAY/SALARY          NaN    Kristin    2271.79    ...

```

```

      transaction_id    country    customer_id \
0 9ca281650e5d482d9e53f85e959baa66    Australia    CUS-1462656821
1 1822eb0e1bbe4c9e95ebbb0fa2cc4323    Australia    CUS-2500783281
2 bd62b1799a454cedbbb56364f7c40cbf    Australia    CUS-326006476
3 0d95c7c932bb48e5b44c2637bdd3efe9    Australia    CUS-1433879684

```

4 f50ccf1195214d14a0acbfc5a265193 Australia CUS-4123612273

	merchant_long_lat	movement	Annual_Salary	age_y	gender_y	amount_y	\
0	NaN	credit	15615.80	23	M	7183.55	
1	NaN	credit	6505.92	23	F	3637.50	
2	NaN	credit	3933.44	22	M	1920.93	
3	NaN	credit	5632.32	24	M	5706.80	
4	NaN	credit	4272.16	43	F	7359.56	

	Annual_Expense
0	28734.0
1	14550.0
2	7684.0
3	22827.0
4	29438.0

[5 rows x 27 columns]

1.2.2 Merging Expense to data

```
[72]: Salary=pd.merge(Salary,Expense,on='customer_id', how='left')
Salary.drop_duplicates(inplace=True)
Salary.head()
```

	index	status	card_present_flag	account	currency	long_lat	\
0	50	posted	NaN	ACC-588564840	AUD	151.27 -33.76	
1	61	posted	NaN	ACC-1650504218	AUD	145.01 -37.93	
2	64	posted	NaN	ACC-3326339947	AUD	151.18 -33.80	
3	68	posted	NaN	ACC-3541460373	AUD	145.00 -37.83	
4	70	posted	NaN	ACC-2776252858	AUD	144.95 -37.76	

	txn_description	merchant_id	first_name	balance	...	customer_id	\
0	PAY/SALARY	NaN	Isaiah	8342.11	...	CUS-1462656821	
1	PAY/SALARY	NaN	Marissa	2040.58	...	CUS-2500783281	
2	PAY/SALARY	NaN	Eric	3158.51	...	CUS-326006476	
3	PAY/SALARY	NaN	Jeffrey	2517.66	...	CUS-1433879684	
4	PAY/SALARY	NaN	Kristin	2271.79	...	CUS-4123612273	

	merchant_long_lat	movement	Annual_Salary	age_y	gender_y	amount_y	\
0	NaN	credit	15615.80	23	M	7183.55	
1	NaN	credit	6505.92	23	F	3637.50	
2	NaN	credit	3933.44	22	M	1920.93	
3	NaN	credit	5632.32	24	M	5706.80	
4	NaN	credit	4272.16	43	F	7359.56	

	Annual_Expense_x	amount	Annual_Expense_y
0	28734.0	7183.55	28734.0


```

1      4  14550.093637050359.56  14550.09438.0
2      7684.0 1920.93          7684.0
3      22827.0 5706.80         22827.0
[5 rows x 29 columns]


```

Scatter plot of annual salary with other features


```

[73]: sns.scatterplot(x= 'age', y= 'Annual_Salary', hue='gender',
                      Anz_dataset =Salary,
                      palette=["purple", "#55CCCC"]);

```



```
[74]: sns.scatterplot(x= 'age', y= 'Annual_balance', hue='gender',
                    Anz_dataset=Salary,
                    palette=["purple", "#55CCCC"]);
```



Dropping irrelevant features

```
[75]: Salary.drop(columns=["customer_id", "amount_x", "amount_y", "balance"], inplace=True)
```

Checking Correlation

```
[76]: Salary.corr()
```

```
[76]:
```

	Annual_Salary	age	Annual_balance	Annual_Expense
Annual_Salary	1.000000	-0.036504	0.217715	0.373476
age	-0.036504	1.000000	0.289224	-0.086174
Annual_balance	0.217715	0.289224	1.000000	0.004943
Annual_Expense	0.373476	-0.086174	0.004943	1.000000

The data looks good with low correlation.

```
[77]: Salary.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 100 entries, 0 to 99
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Annual_Salary    100 non-null    float64
1   age              100 non-null    int64
2   gender           100 non-null    object
3   Annual_balance   100 non-null    float64
4   Annual_Expense   100 non-null    float64
dtypes: float64(3), int64(1), object(1)
memory usage: 4.7+ KB
```

One Hot Encoding of gender

```
[78]: Salary = pd.get_dummies(Salary, columns=["gender"])
```

```
[79]: Salary.head()
```

```
[79]:
```

	Annual_Salary	age	Annual_balance	Annual_Expense	gender_F	gender_M
0	50464.44	53	245370.60	14611.0	1	0
1	100202.20	21	334801.68	35735.0	0	1

2	45996.24	28	140201.28	22046.0	0	1
3	88992.28	34	458300.32	26931.0	1	0
4	109304.44	34	667680.08	34898.0	1	0

[80]: Salary.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 100 entries, 0 to 99
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Annual_Salary    100 non-null    float64
1   age              100 non-null    int64
2   Annual_balance   100 non-null    float64
3   Annual_Expense   100 non-null    float64
4   gender_F         100 non-null    uint8
5   gender_M         100 non-null    uint8
dtypes: float64(3), int64(1), uint8(2)
memory usage: 4.1 KB
```

Linear regression Model

[81]: `x = Salary.drop("Annual_Salary", axis="columns")`
`y = Salary["Annual_Salary"]`

[82]: *#initialize the linear regression model*
`reg = linear_model.LinearRegression()`

[83]: *#split the data into 70% training and 33% testing data*
`x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.3,random_state=7)`

[84]: *#train the model*
`reg.fit(x_train, y_train)`

[84]: LinearRegression()

[85]: *#print the coefficients/weights for each feature/column of our model*
`print(reg.coef_)`

```
[-2.19854918e+02  2.98116645e-03  8.77244212e-01 -3.24159489e+03
  3.24159489e+03]
```

[86]: `y_pred = reg.predict(x_test)`
`print(y_pred)`

```
[66169.19169742  73238.40211721  67870.27501126  82291.76186935
 58072.28717096  57881.30249663  69273.84842515  57906.52412454]
```

```
67626.99063658 75146.47780116 77691.48793785 64701.99935058
65233.22702201 71287.20044804 62880.56730571 62093.20570672
56473.86099523 61340.1809038 75227.34041434 68533.01201909
71981.9013954 69816.96105331 78150.49074061 54270.48300664
61215.47635446 75355.5409963 54143.23057579 73745.14836422
76595.75997825 64676.66144645]
```

```
[87]: #print the actual values
      print(y_test)
```

```
37    127048.48
26     85991.92
78     36904.32
91    118049.12
49     39128.64
15     59290.80
93     85109.44
71     44235.36
86     51508.60
22     61538.96
13    109310.60
40     53249.52
52     44873.40
12     39426.24
88     39287.20
45     39379.92
11     58414.72
66     44004.00
20     78147.44
18     46388.68
50    123348.40
2      45996.24
17     68513.76
85     44550.72
5      63906.08
97     43406.88
51     44194.08
30     55408.08
74     55111.68
96     53927.64
```

Name: Annual_Salary, dtype: float64

```
[88]: reg.score(x_train,y_train)
```

```
[88]: 0.1741029921866517
```

```
[89]: fig, ax = plt.subplots()
      ax.scatter(y_test, y_pred, edgecolors=(0, 0, 0))
```

```
ax.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--', lw=4)
ax.set_xlabel('Actual')
ax.set_ylabel('Predicted')
ax.set_title("Annual Salary vs Predicted")
plt.show()
```



The linear regression model accuracy was 17.4%

DECISION TREE

```
[90]: x = Salary.drop('Annual_Salary', axis='columns')
      y = Salary['Annual_Salary']
```

```
[91]: x_train, x_test, y_train, y_test = train_test_split(x,y,
      test_size=0.3,random_state=7)
```

```
[92]: dt = DecisionTreeRegressor(max_depth=4,min_samples_leaf=0.1,random_state=7)
```

```
[93]: dt.fit(x_train, y_train)
```

```
[93]: DecisionTreeRegressor(max_depth=4, min_samples_leaf=0.1, random_state=7)
```

```
[94]: y_pred = dt.predict(x_test)
      print(y_pred)
```

```
[ 76408.79428571  96712.74285714  52270.156      96712.74285714
  52270.156      96712.74285714  76408.79428571  44048.15076923
 64360.90153846  76408.79428571  76408.79428571  44048.15076923
 70733.33333333  52270.156      52270.156      44048.15076923
 64360.90153846  44048.15076923  106748.31      70733.33333333
 76408.79428571  52270.156      76408.79428571  44048.15076923
 70733.33333333  52270.156      64360.90153846  70733.33333333
 70733.33333333  70733.33333333]
```

```
[95]: print(y_test)
```

```
37    127048.48
26     85991.92
78     36904.32
91    118049.12
49     39128.64
15     59290.80
93     85109.44
71     44235.36
86     51508.60
22     61538.96
13    109310.60
40     53249.52
52     44873.40
12     39426.24
88     39287.20
45     39379.92
11     58414.72
66     44004.00
20     78147.44
18     46388.68
50    123348.40
2      45996.24
17     68513.76
85     44550.72
5      63906.08
97     43406.88
51     44194.08
30     55408.08
74     55111.68
96     53927.64
```

Name: Annual_Salary, dtype: float64

```
[96]: dt.score(x_train,y_train)
```

```
[96]: 0.5600863941994056
```

```
[97]: mse_dt = MSE(y_test, y_pred)
```

```
[98]: rmse_dt = mse_dt**(1 / 2)
```

```
[99]: print(rmse_dt)
```

20527.660154490288

Decision Tree Regression accuracy was 56%.

Conclusion

The decision was the preferred algorithm for now because it more efficient than linear regression. Linear regression accuracy was 17 percent while that of decision tree was 56 percent. We may use random forest model or XGBoost Model subsequently to see their performance in predicting customers' incomes bracket which is an important feature in this case,